

# Deep Learning for Computer Vision

Lecture 2: An Introduction to Old-School Machine Learning

Peter Belhumeur

Computer Science  
Columbia University

# Who are these people?

---



[Sinha and Poggio 1996]

# Who are these people?


---



[Sinha and Poggio 2002]

# Images as Points in Euclidean Space


---



- Let an  $n$ -pixel image to be a point in an  $n$ -D space,  $\mathbf{x} \in \mathbb{R}^n$ .
- Each pixel value is a coordinate of  $\mathbf{x}$ .

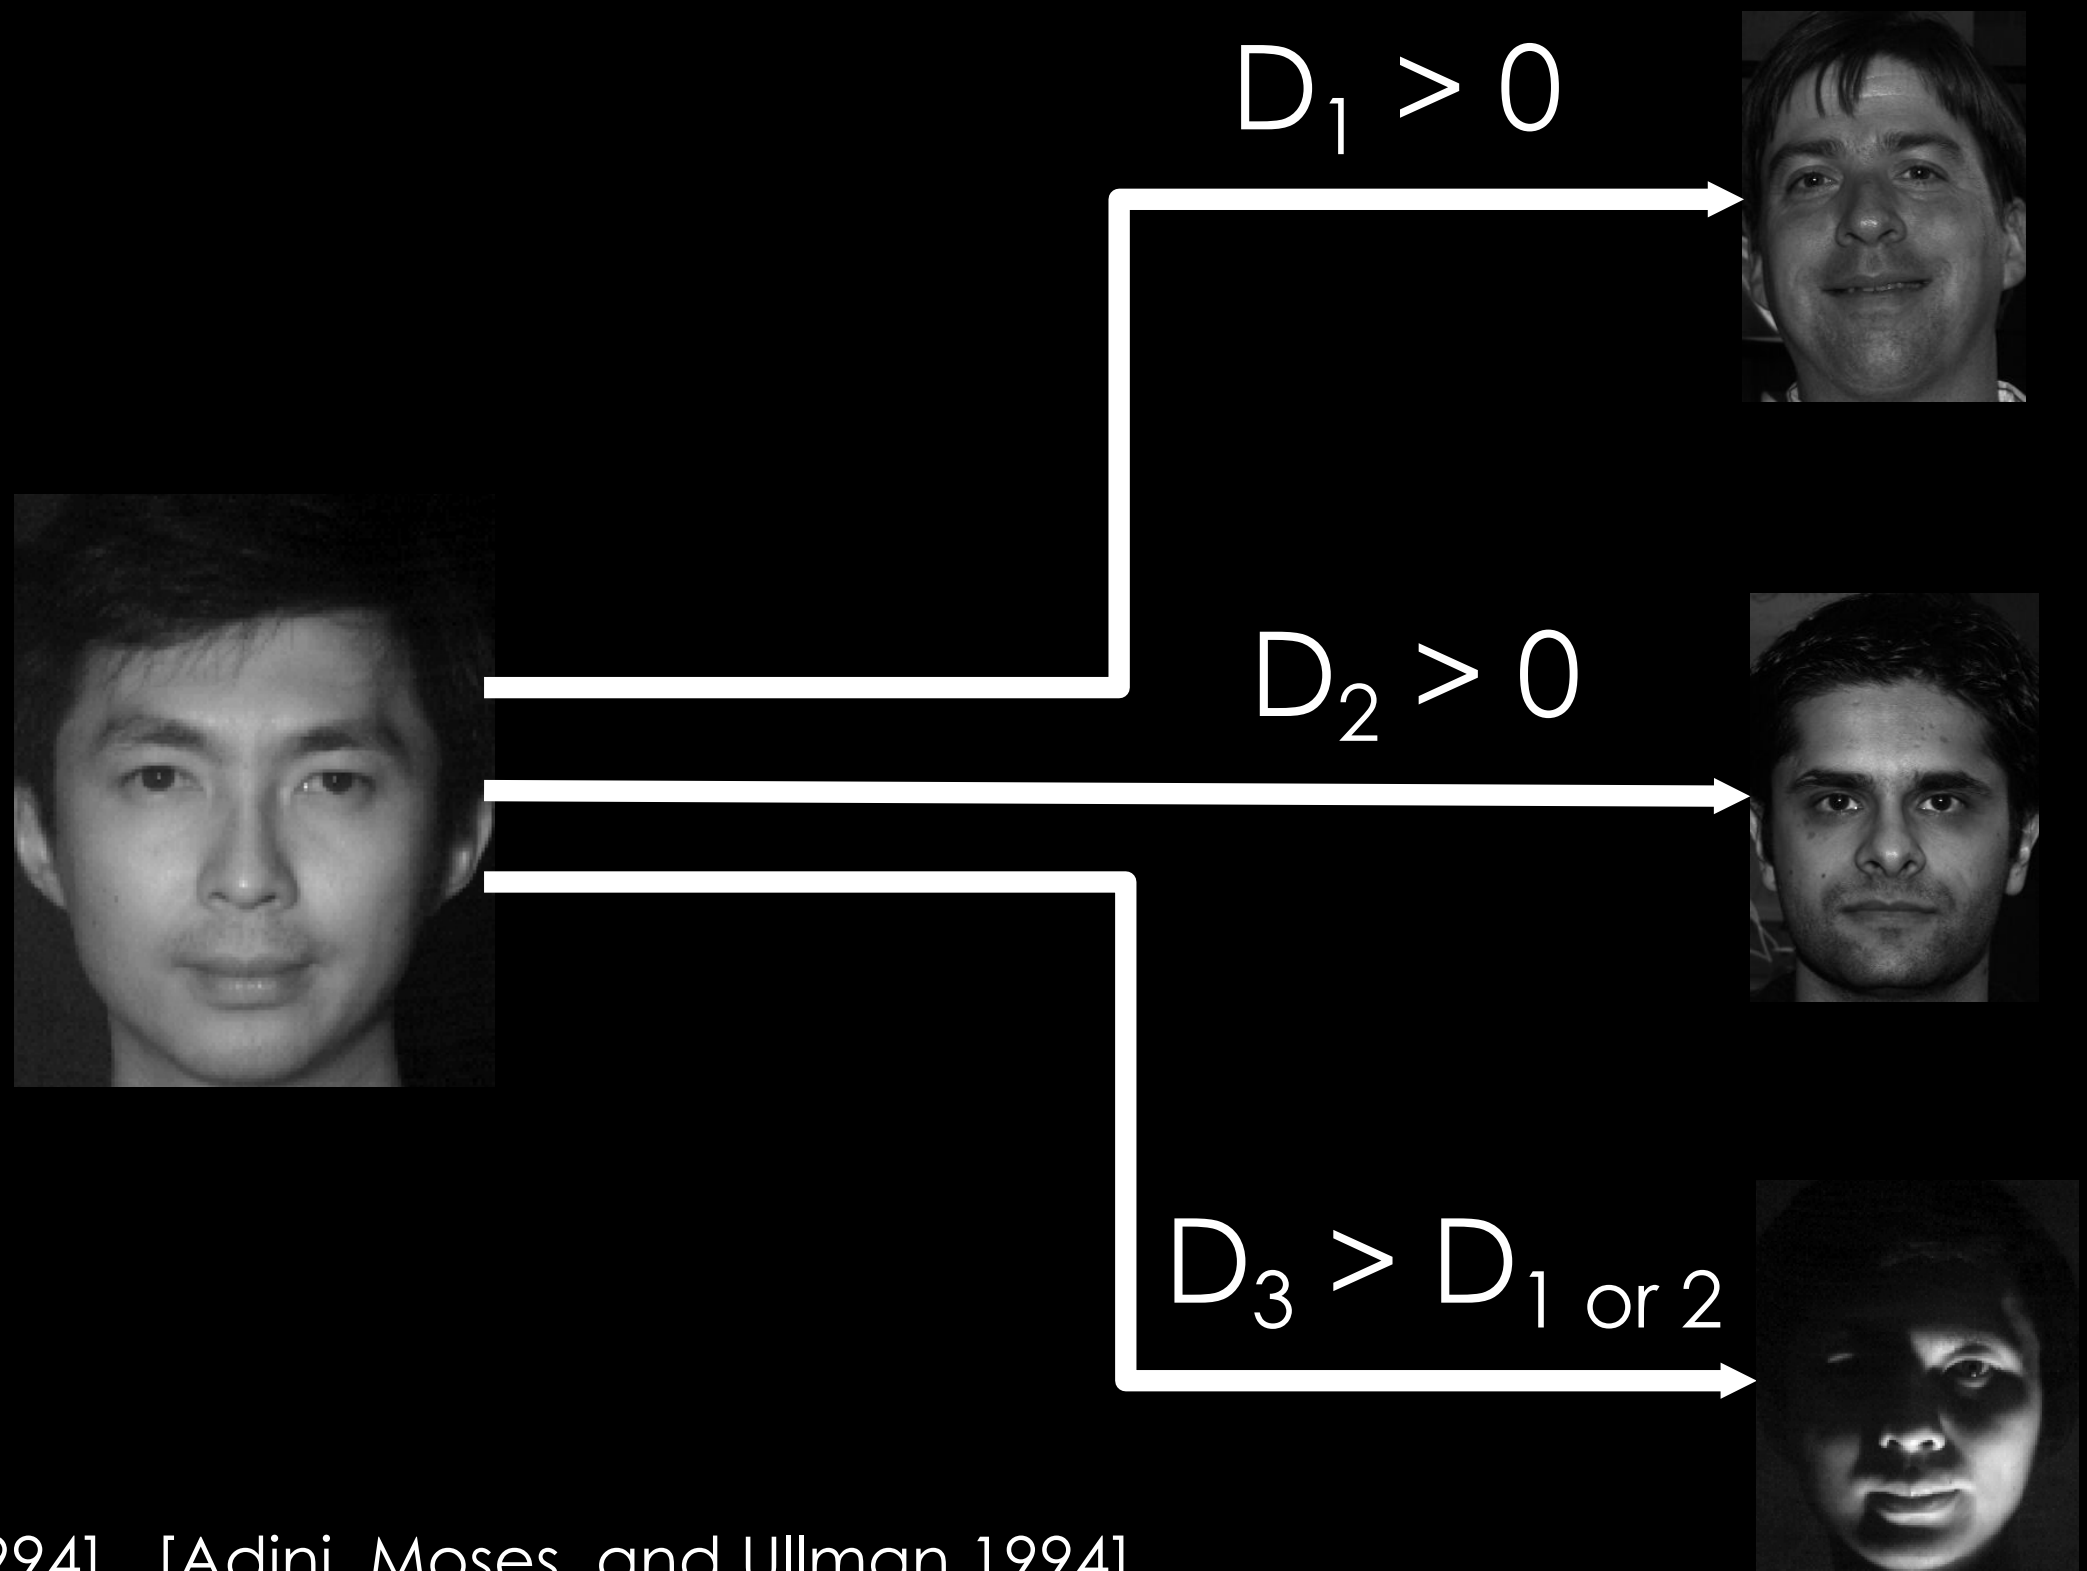
# Image Recognition: Euclidean Distances

---



# Image Recognition: Euclidean Distances

---



**Same Person  
or  
Different People**





**Same Person  
or  
Different People**







# Non-Existence Theorem for Geometric Invariants

---

Geometric invariants for rigid transformations of 3-D objects viewed under perspective projective projection **do not exist**.


# Non-Existence Theorem for Illumination Invariants

---

Illumination invariants for 3-D objects **do not exist**.

This result **does not ignore** attached and cast shadows, as well as surface interreflection.

# Classifiers




# Pattern Classification

Duda, Hart, and Stork

# Classification: An Example

Classify fish species at an Alaskan Canning Factory

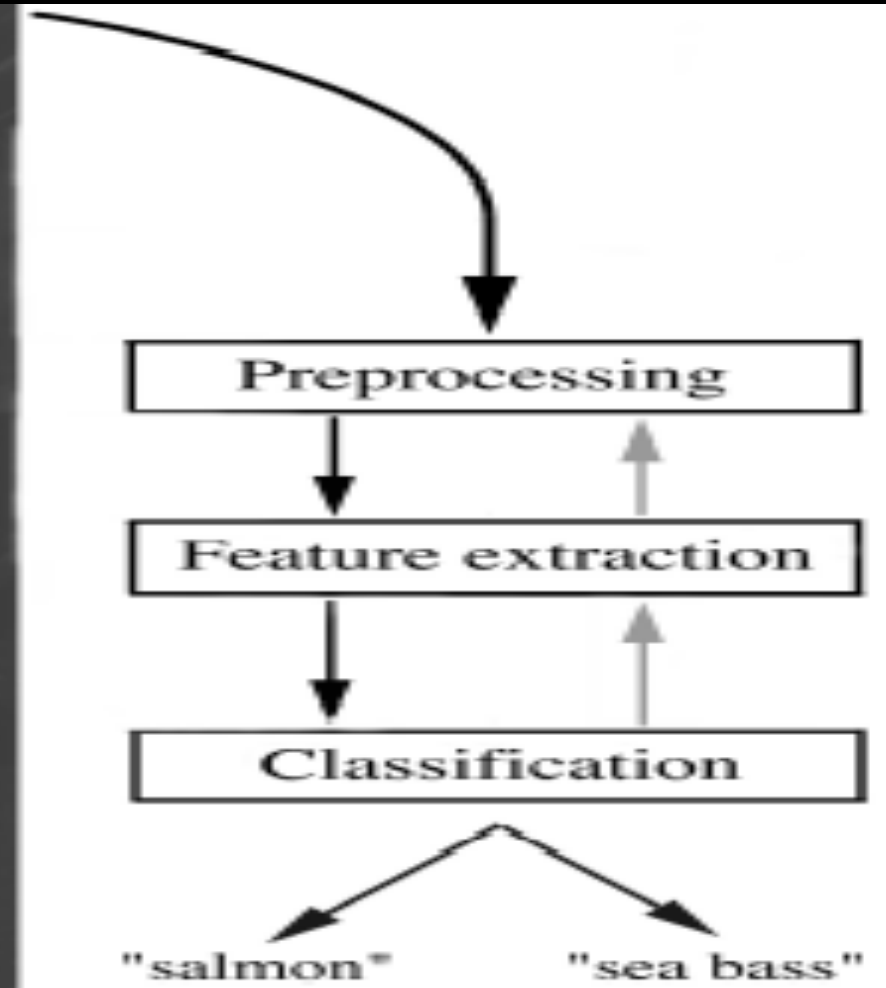




# Building a classifier


Use a camera and take some sample images to extract features

- Length
  - Lightness
  - Width
  - Number and shape of fins
  - Position of the mouth, etc...
- 
- This is the set of all suggested features to explore for use in our classifier!



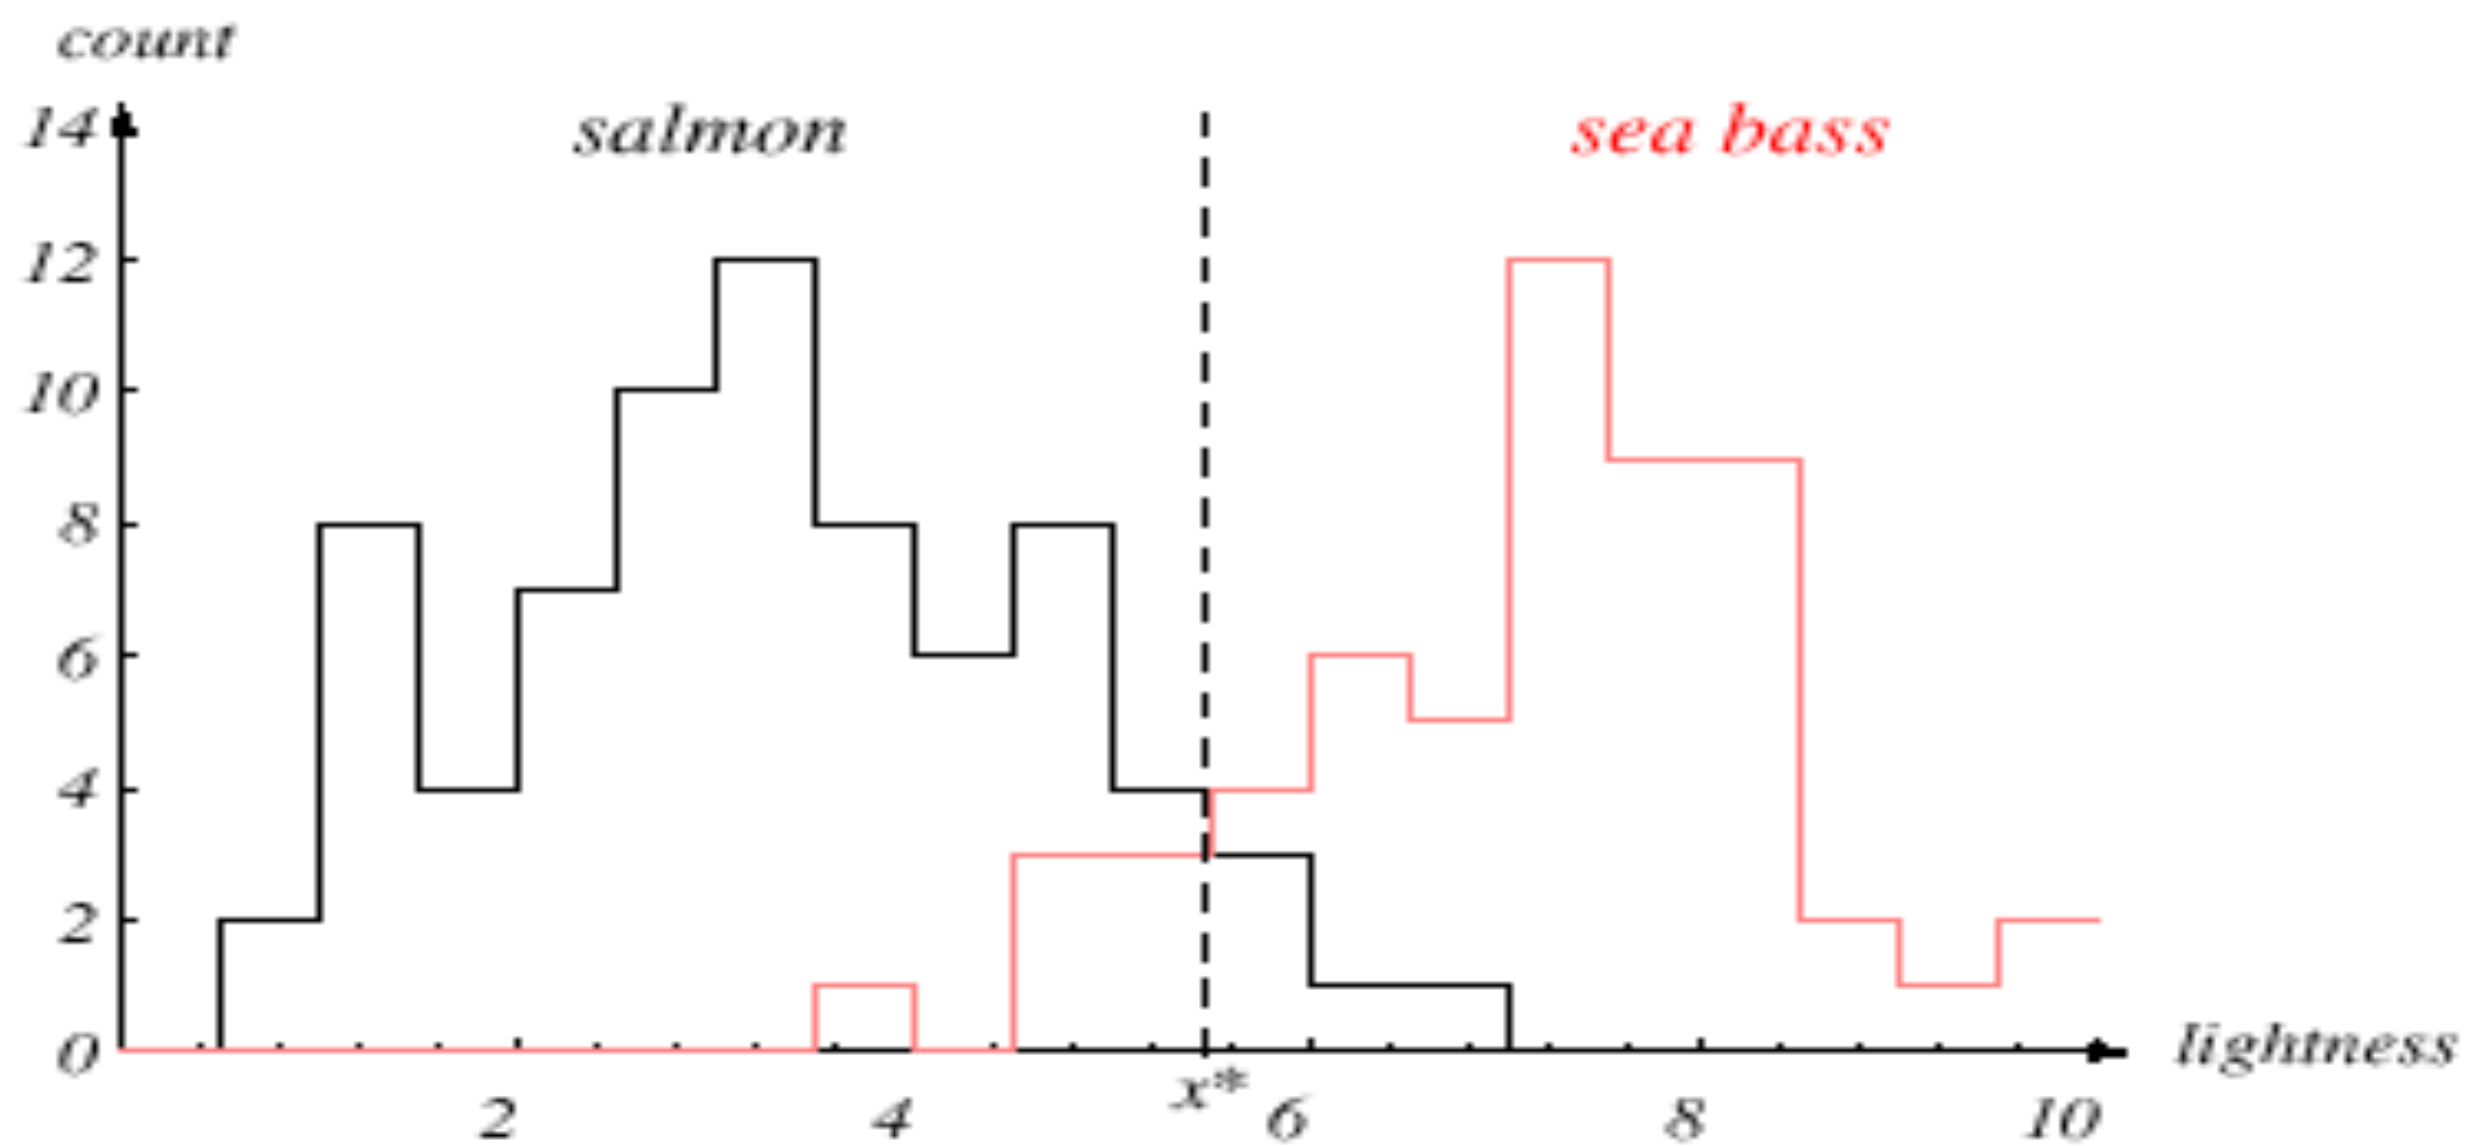
# Manual Feature Selection

Select the length of the fish as a possible feature for discrimination




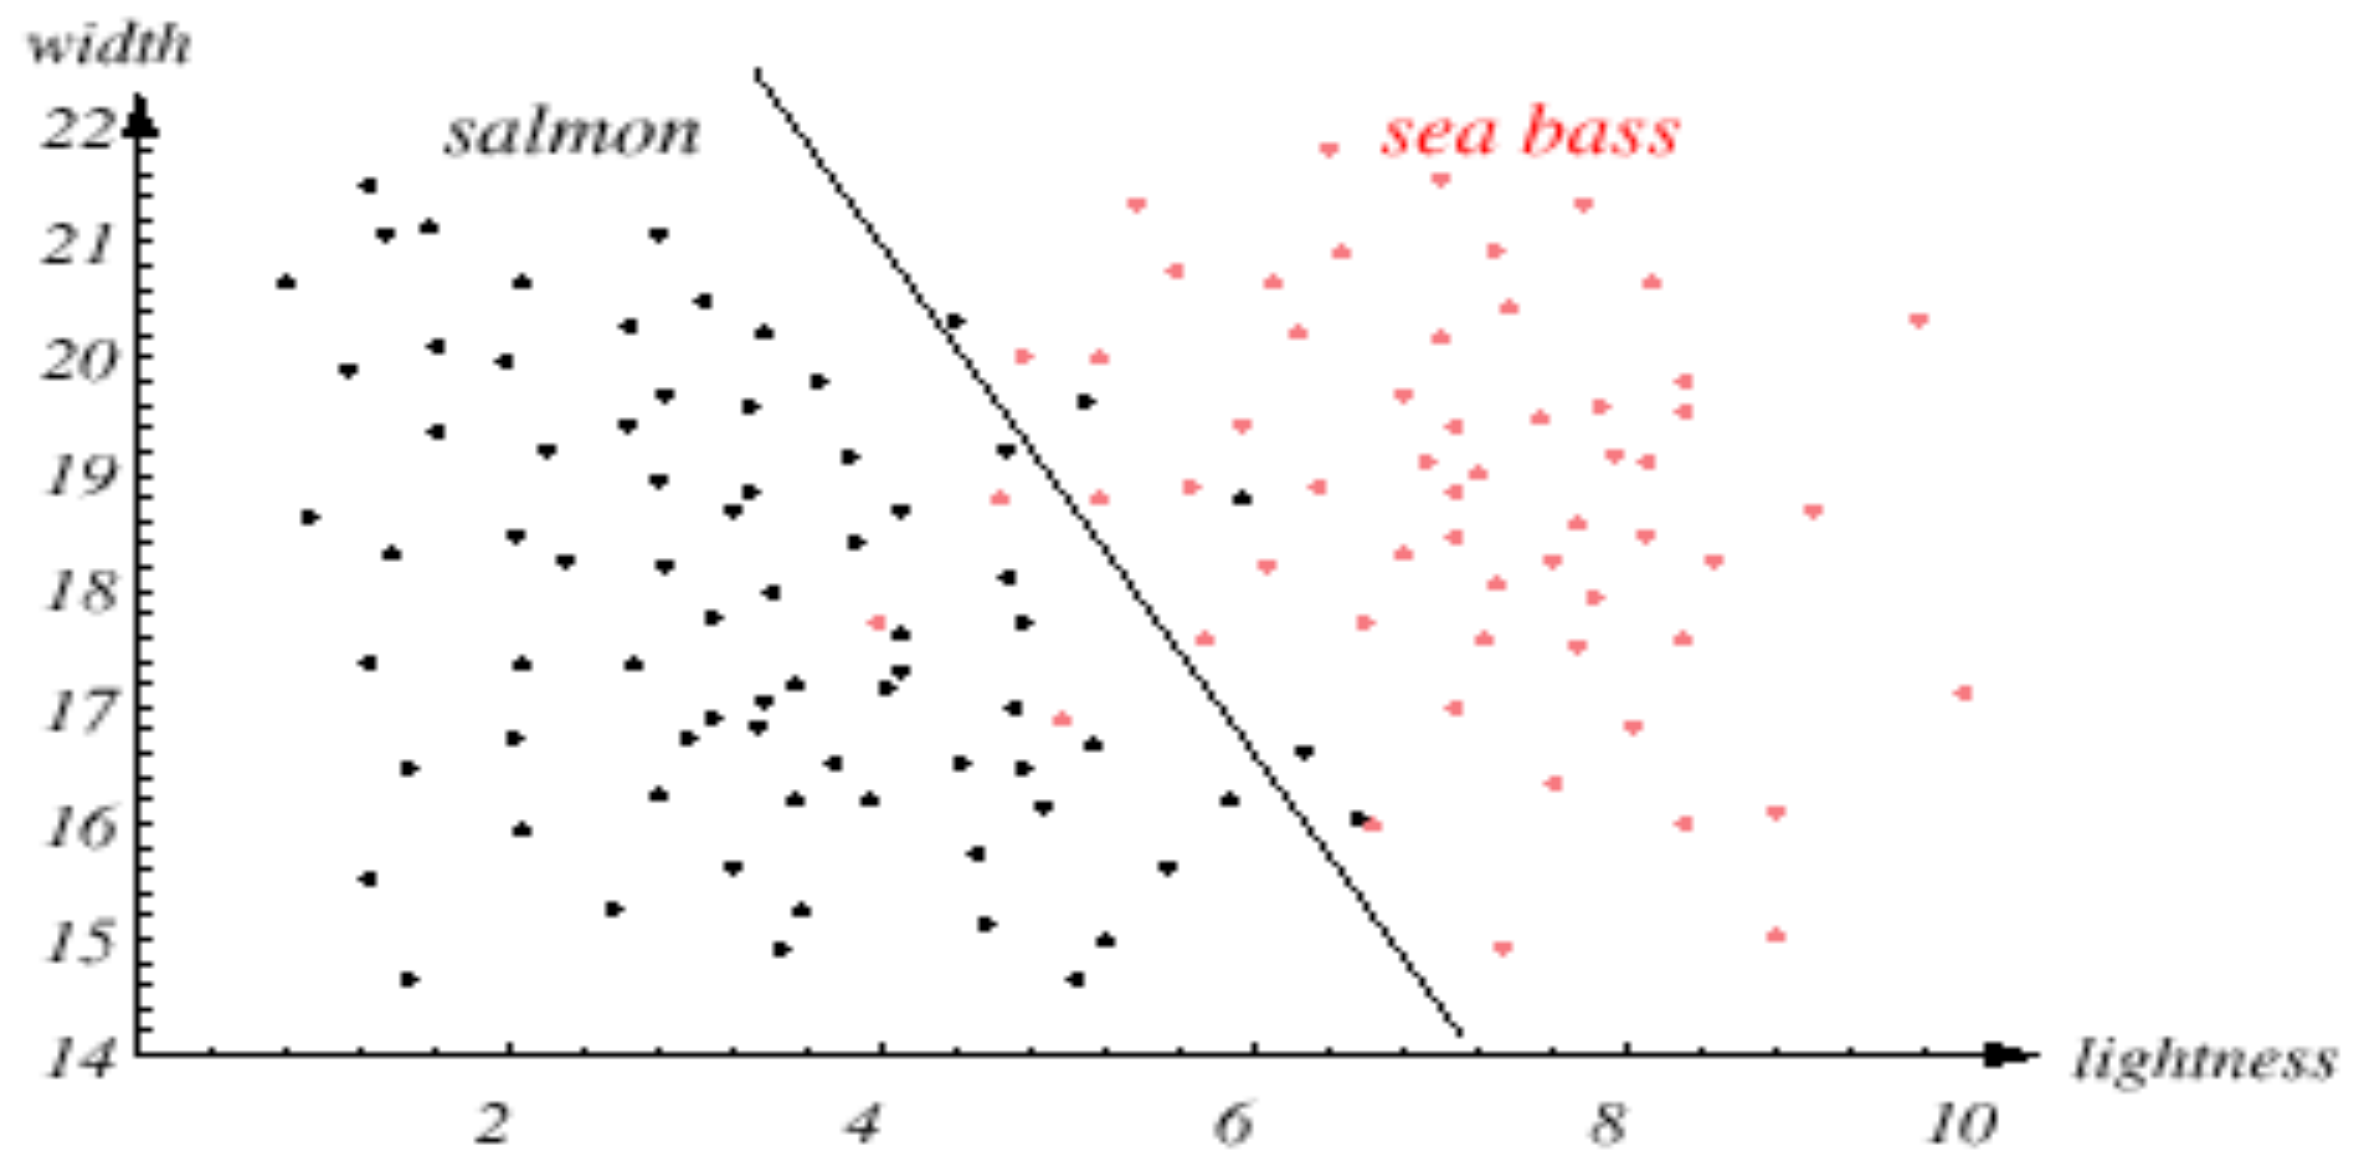
The length is a **poor** feature alone!

Select the lightness as a possible feature.



Adopt the lightness and add the width of the fish









We might add other features that are not correlated with the ones we already have. A precaution should be taken not to reduce the performance by adding “noisy features”

Ideally, the best decision boundary should be the one which provides an optimal performance.



However, our satisfaction is premature because the central aim of designing a classifier is to correctly classify novel input



Issue of generalization!

